

# Introduction to AI

Assignment Project Exam Help

-Search-

<https://tutorcs.com>

Francesca Toni

WeChat: cstutorcs

# Outline


- Solving problems by search
- Basic search algorithms (uninformed/blind search)
- Heuristic-based search algorithms (informed search)
- Search for games (adversarial search )

Assignment Project Exam Help


<https://tutorcs.com>

WeChat: cstutorcs

Recommended reading:  
(most of) Chapters 3, 5



Additional reading:  
Chapter 3  
Section 11.2.2



# Solving problems by search

- Problem solution can be abstracted as **path** from some (**start**) node to some (**goal**) node in a suitably defined **directed graph**
- Search algorithms to find (**optimal**) paths

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Route finding example


On holiday in Romania; today in Arad; flight leaves tomorrow from Bucharest

- Start: Arad
- Goal: Bucharest
- Optimal path
  - shortest, or
  - quickest, or
  - cheapest

Assignment Project • Exam Help


<https://tutorcs.com>

WeChat: cstutorcs



# Vacuum world example

The vacuum cleaner needs to remove all dirt from the room



- Start: A


Assignment Project Exam Help

- Goal: no dirt anywhere in A, B


<https://tutorcs.com>

WeChat: cstutorcs

- Optimal path
  - shortest, or
  - quickest, or
  - cheapest



# 8-puzzle example



- Optimal path?
- Graph?

# Which problems can be solved by (classical) search?

Environment is

- Observable
  - Current state known
- Discrete
  - From each state finitely many next states (by executing actions)
- Deterministic
  - Each action has exactly one outcome (next state)
- Known
  - Possible actions and next states for each state

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Search problems – formally


- *Initial state (start)* e.g.  $In(Arad)$
- *Transition model* between states (**graph**)  
possibly given by *successor function*  $RESULT: States \times Actions \rightarrow States$   
e.g.  $RESULT(In(Arad), do(Zerind)) = In(Zerind)$
- *Goal test* (set of **goal** states)  
<https://tutorcs.com>  
• explicit e.g.  $In(Bucharest)$ , or  
• implicit WeChat: [cstutorcs](https://tutorcs.com) e.g.  $\forall x \neg Dirt(x)$
- *Path cost* (additive = sum of positive step costs )  
e.g. sum of distances, number of actions

**Solution**=path (sequence of actions) from the initial state to a goal state


**Optimal solution**= solution with lowest path cost



# Vacuum world example – formally



- Initial state=  $In(A) \wedge Dirt(A) \wedge Dirt(B)$
- Goal test=  $\{\neg Dirt(A) \wedge \neg Dirt(B)\}$
- RESULT ( $In(A) \wedge Dirt(A) \wedge Dirt(B), S$ ) =  $In(A) \wedge \neg Dirt(A) \wedge Dirt(B), etc$



- Step cost=1, path cost=number of actions

# 8-puzzle example – formally

7	2	4
5		6
8	3	1

1	2	3
4	5	6
7	8	

Start

Goal

<https://tutorcs.com>

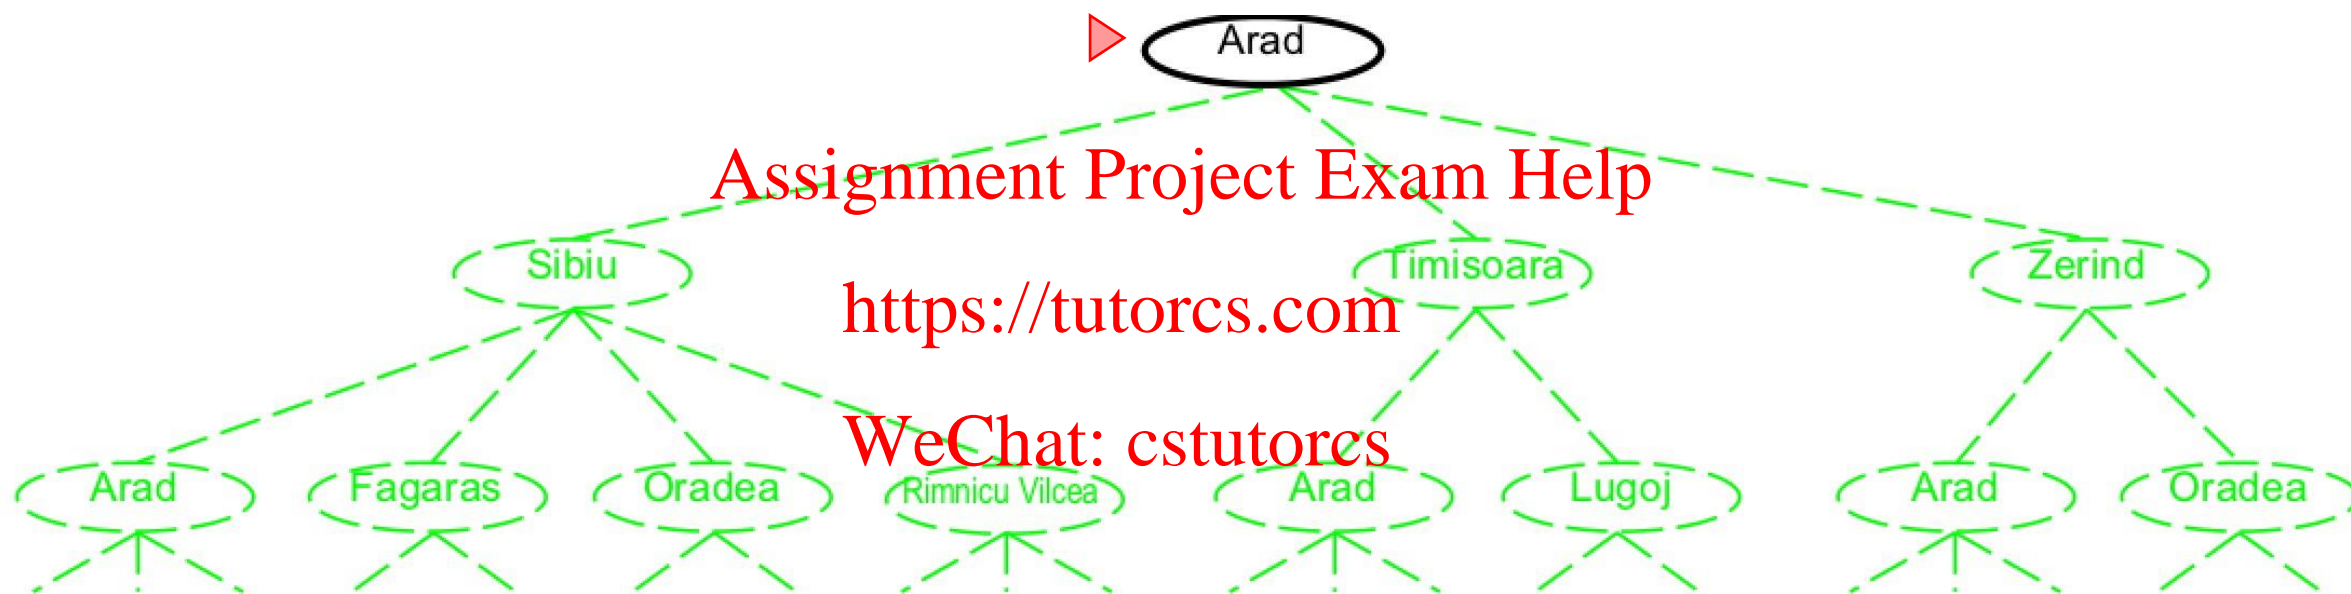
WeChat: cstutorcs

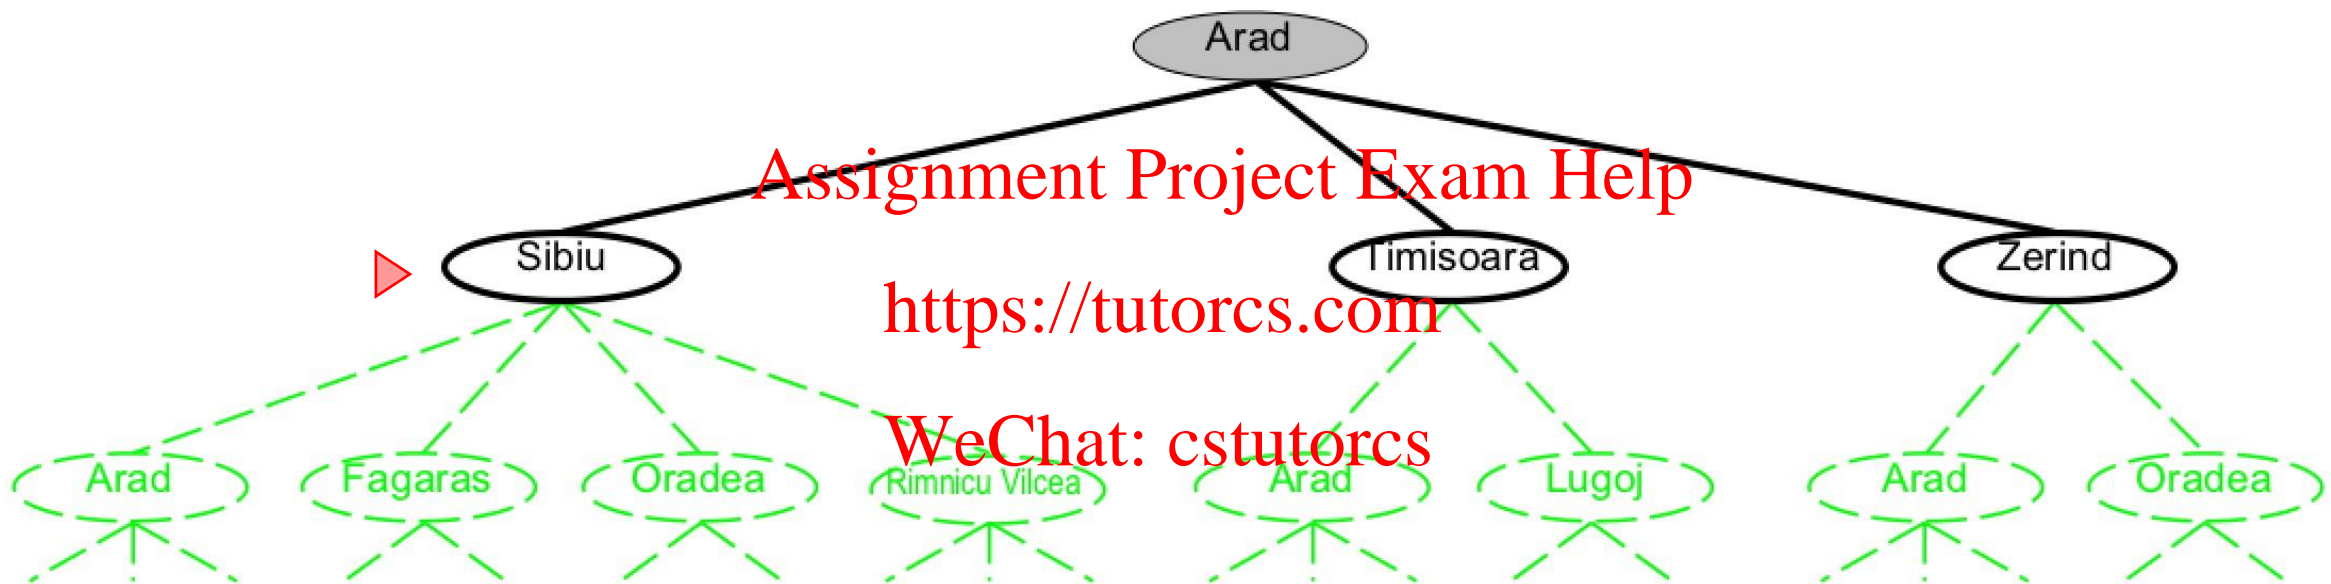
- *Initial state* = Start – e.g. represented by integer locations of tiles
- *Transition model* between states?  
(*successor function* RESULT: States×Actions → States ?)
- *Goal test* = {Goal} – e.g. represented by integer locations of tiles
- *Path cost* = number of actions

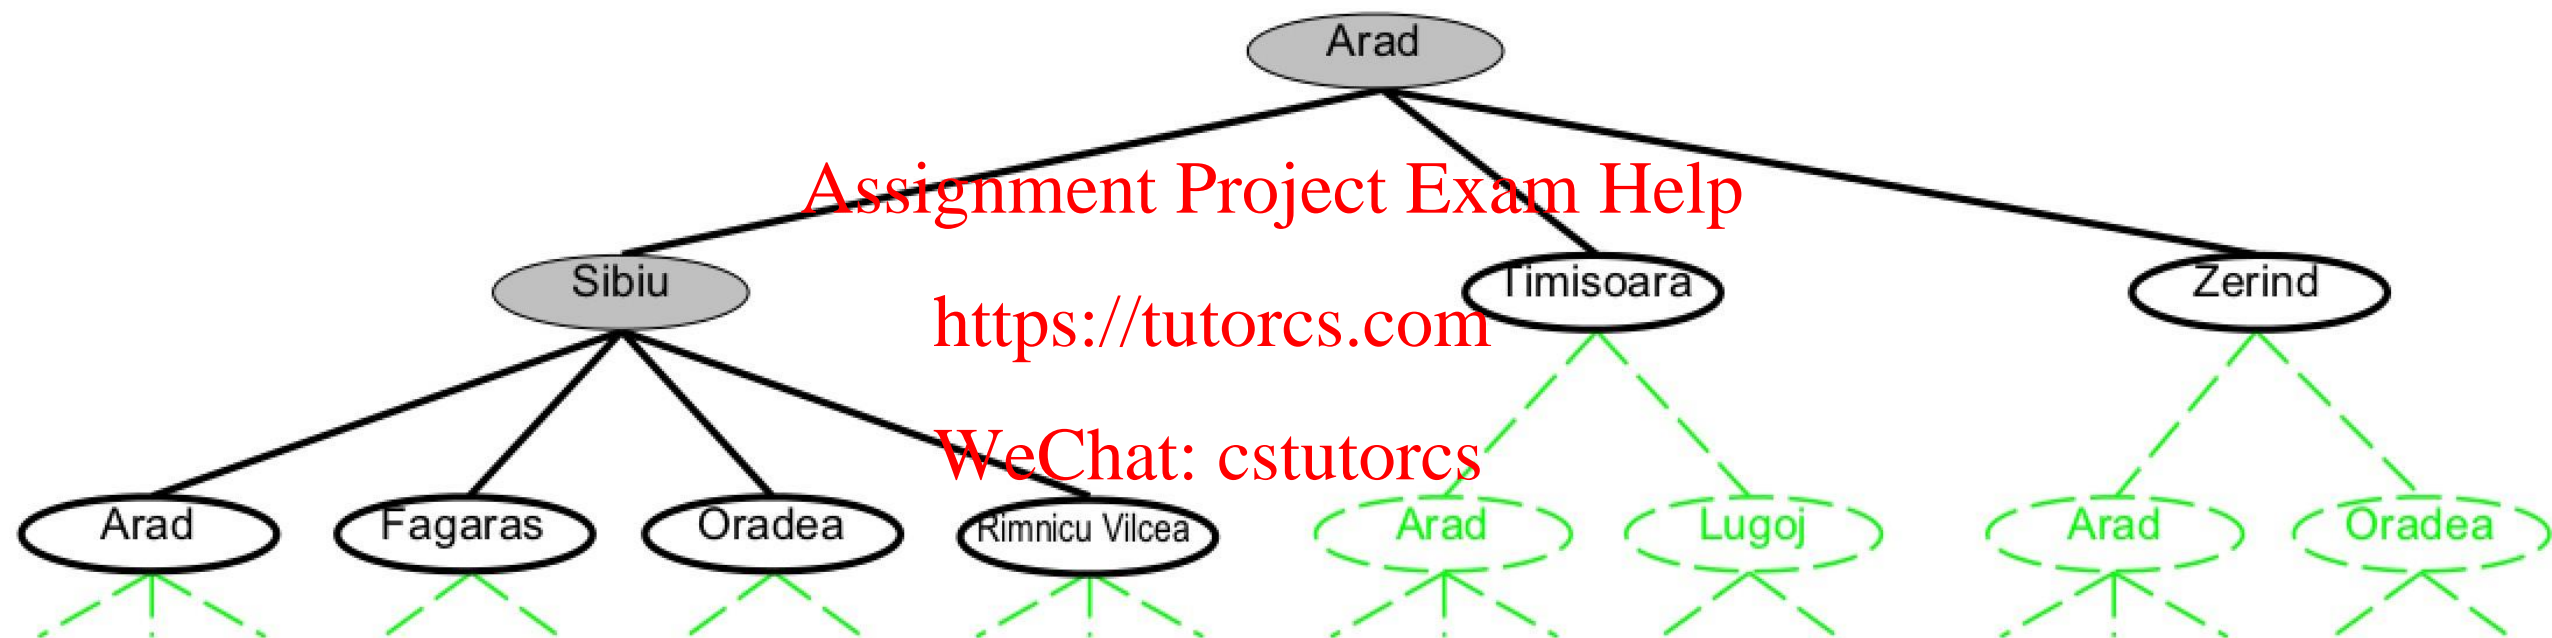
# Search algorithms – basic idea

Generate a (search) tree:

1. Initialise the tree with the initial state as root (=frontier)
2. Repeat
  - i. if the frontier of the tree is empty then return fail
  - ii. **choose** and remove a leaf node L from the frontier
  - iii. if L is a goal state (in the goal test) then return corresponding solution
  - iv. expand L, adding the resulting nodes to the frontier







# Search algorithms – basic idea with loop checking

Generate a (search) tree:

1. Initialise the tree with the initial state as root
2. Repeat
  - i. if the frontier of the tree is empty then return fail
  - ii. **choose** and remove a leaf node L from the frontier
  - iii. if L is a goal state (in the goal test) then return corresponding solution
    - \*add L to the “explored set”
  - iv. expand L, adding the resulting nodes to the frontier
    - \*only if not in the frontier or the “explored set” already

\* Additions to avoid loops

# Search algorithms - concretely

Assignment Project Exam Help

<https://tutorcs.com>  
**Definition of choose**

WeChat: cstutorcs



# Uninformed/blind search


- **Breadth-first** search
  - **choose** (a) **shallowest** (closest to root) unexpanded node
- **Uniform-cost** search
  - **choose** unexpanded node with **lowest path cost**
- **Depth-first** search
  - **choose** **deepest** (farthest from the root) unexpanded node.
- **Depth-limited** search
  - depth-first search with given depth limit  $\ell$  (nodes of depth  $\ell$  have no children)
- **Iterative deepening** search
  - depth-limited search of increasing  $\ell$

Assignment Project Exam Help

<https://tutorcs.com>


WeChat: cstutorcs

- Breadth-first (**choose** shallowest)

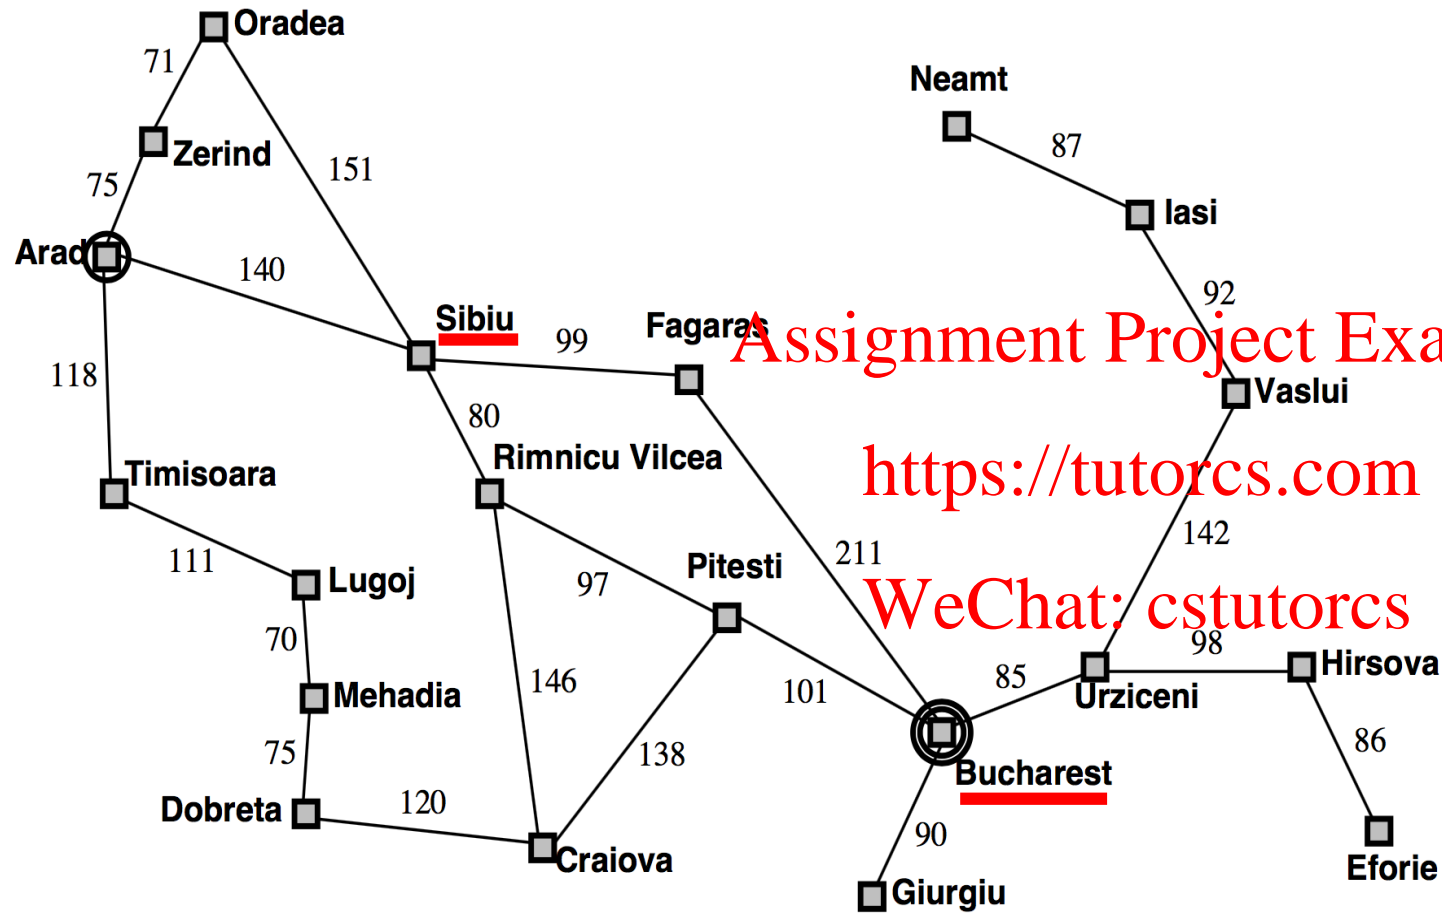


<https://tutorcs.com>

- Uniform-cost (**choose** cheapest) all costs equal = Breadth-first



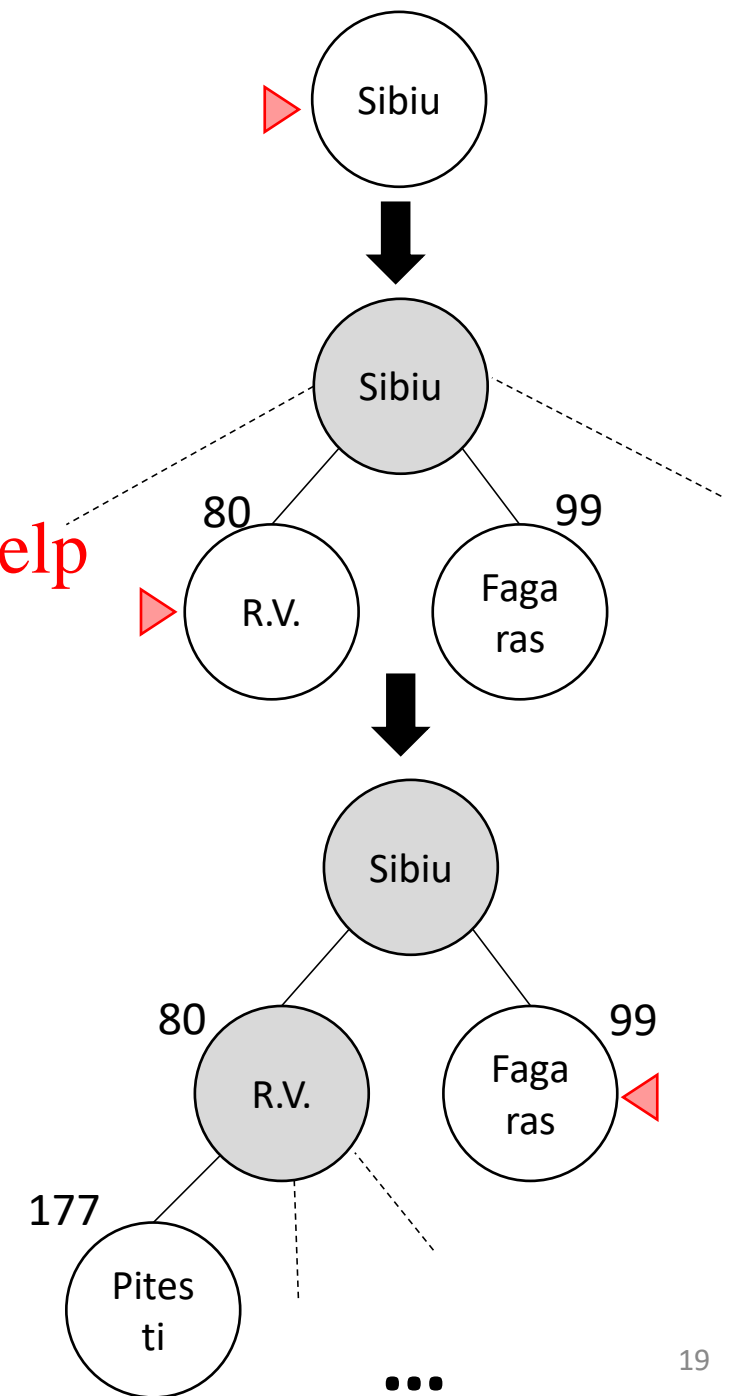
- Uniform-cost (**choose** cheapest) any costs



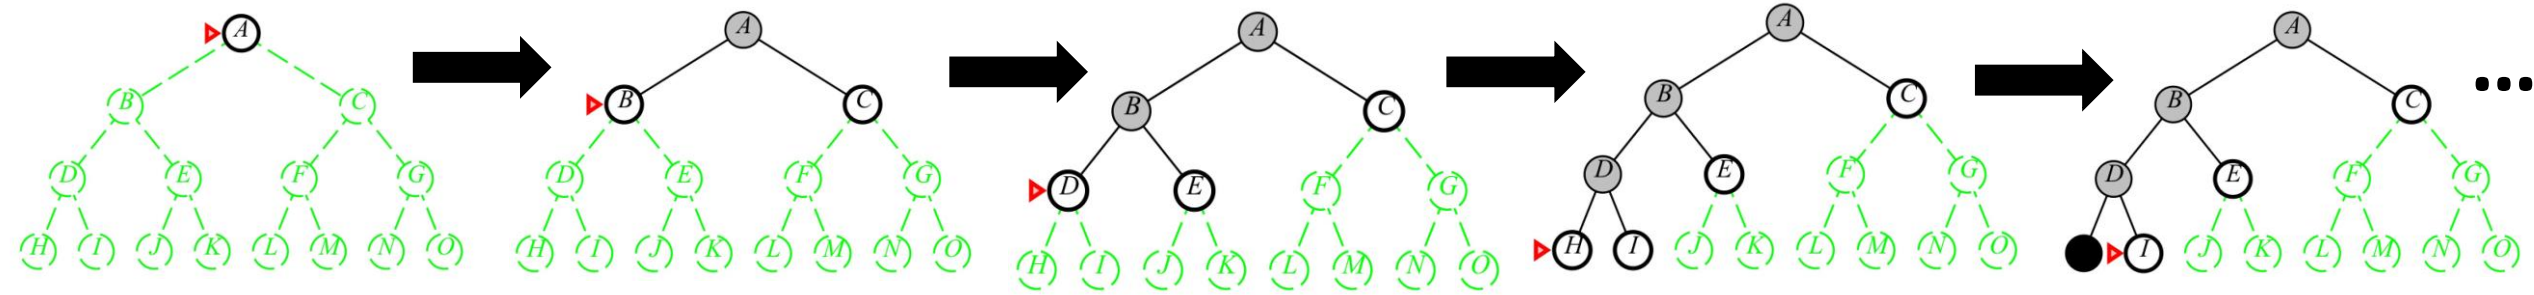
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



- Depth-first (**choose** deepest)




- Depth-limited ( $\ell=1$ )

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



- Iterative deepening

Depth-limited ( $\ell=0$ )  $\longrightarrow$  Depth-limited ( $\ell=1$ )  $\longrightarrow$  Depth-limited ( $\ell=2$ )  $\dots$

# Uninformed search - Variants

- Backtracking search – variant of depth-first search
  - Only one (successor) node generated when expanded
  - Each (partially expanded) successor node remembers which node to generate next
- Bi-directional search
  - Two simultaneous searches (from start to goal and from goal to start)
  - Succeeds when (if) the frontiers of the two searches intersect

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Properties

- *completeness*—does the algorithm always find a solution if one exists?
- *time complexity*—number of nodes generated/expanded
- *space complexity*—maximum number of nodes in memory
- *optimality*—does the algorithm always find an optimal (least-cost) solution?

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Properties of uninformed search (with loop-checking)

Criterion	Breadth-first	Uniform-cost	Depth-first	Backtracking	Depth-limited	Iterative deepening	Bidirectional (breadth-first)
Complete?	Yes*	Yes* <sup>o</sup>	Yes <sup>†</sup>	Yes <sup>†</sup>	No	Yes*	Yes*
Time	$O(b^d)$	$O(b^{d+1})$ ●	$s$ $O(bm)$	$s$ $O(m)$	$s(\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{d+1})$ ●	$O(bm)$	$O(m)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes●	Yes	No	No	No	Yes●	Yes●

\* **b** is finite

● step cost is constant

○ step cost is  $>0$

†  $s$  is finite

**b**—maximum branching factor of the search tree (may be  $\infty$ )

**d**—depth of the optimal (least-cost) solution

**m**—maximum depth of the state space (may be  $\infty$ )

**s**—overall size of the state space

**s( $\ell$ )**—size of the state space till depth  $\ell$

# Properties of uninformed search (without loop-checking)

Criterion	Breadth-first	Uniform-cost	Depth-first	Backtracking	Depth-limited	Iterative deepening	Bi-directional (breadth-first)
Complete?	Yes*	Yes* <sup>o</sup>	No	No	No	Yes*	Yes*
Time	$O(b^d)$	$O(b^{d+1})$ ●	$O(b^m)$	$O(b^m)$	$O(b^l)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{d+1})$ ●	$O(bm)$	$O(m)$	$O(b^l)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes●	Yes	No	No	No	Yes●	Yes●

- \* **b** is finite
- step cost is constant
- step cost is  $>0$

**b**—maximum branching factor of the search tree (may be  $\infty$ )  
**d**—depth of the optimal (least-cost) solution  
**m**—maximum depth of the state space (may be  $\infty$ )



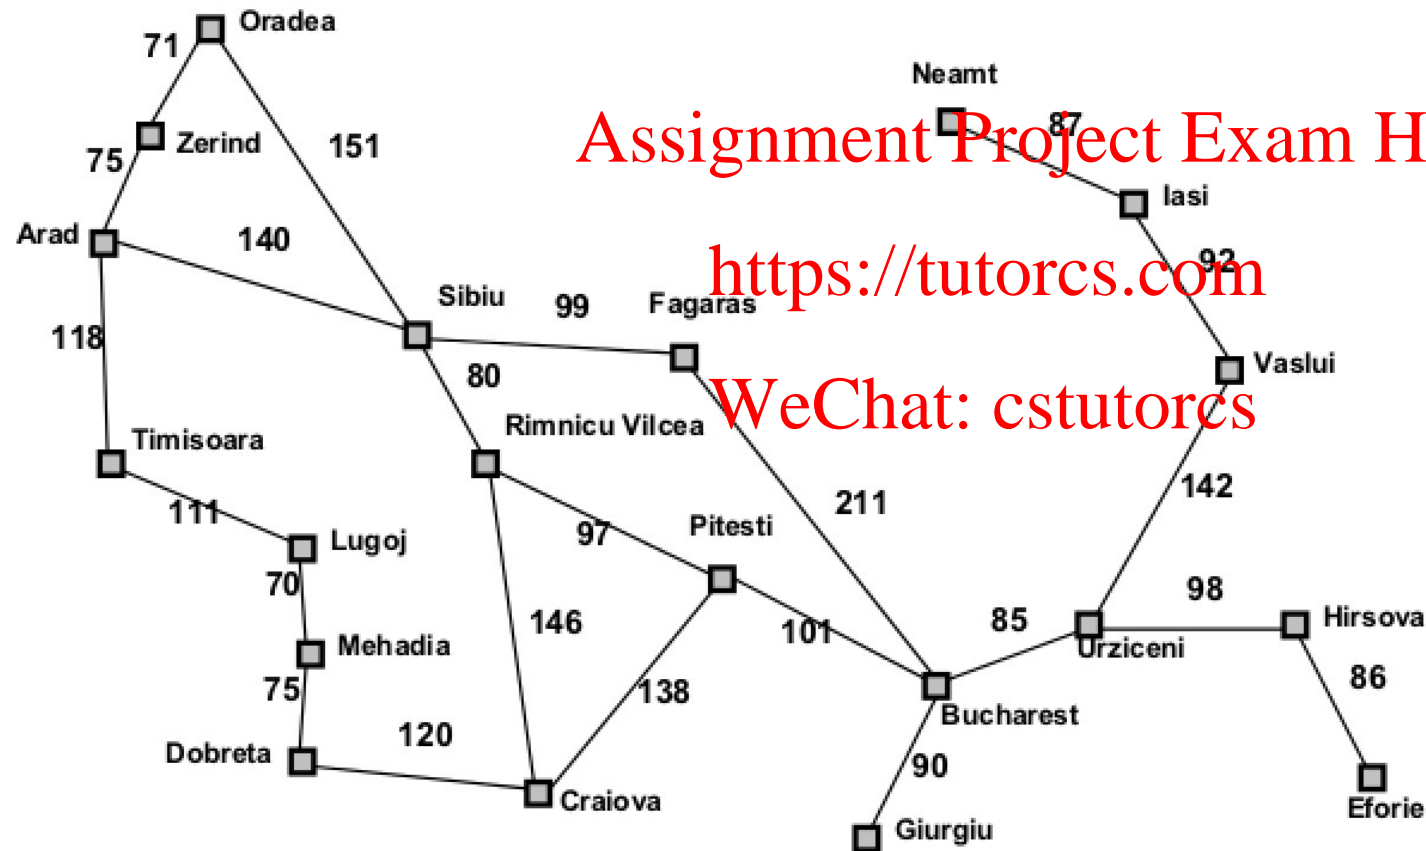
# Informed (heuristic-based) search

Use a *cost estimate* (of optimal path to goal) to **choose** node with the least-estimated path cost

<https://tutorcs.com>

- Greedy-best-first search
  - cost estimate = **heuristic function**(node to goal)
- A\* search
  - cost estimate = **actual cost**(to node) + **heuristic function**(node-to-goal)

# Heuristic function: route finding example

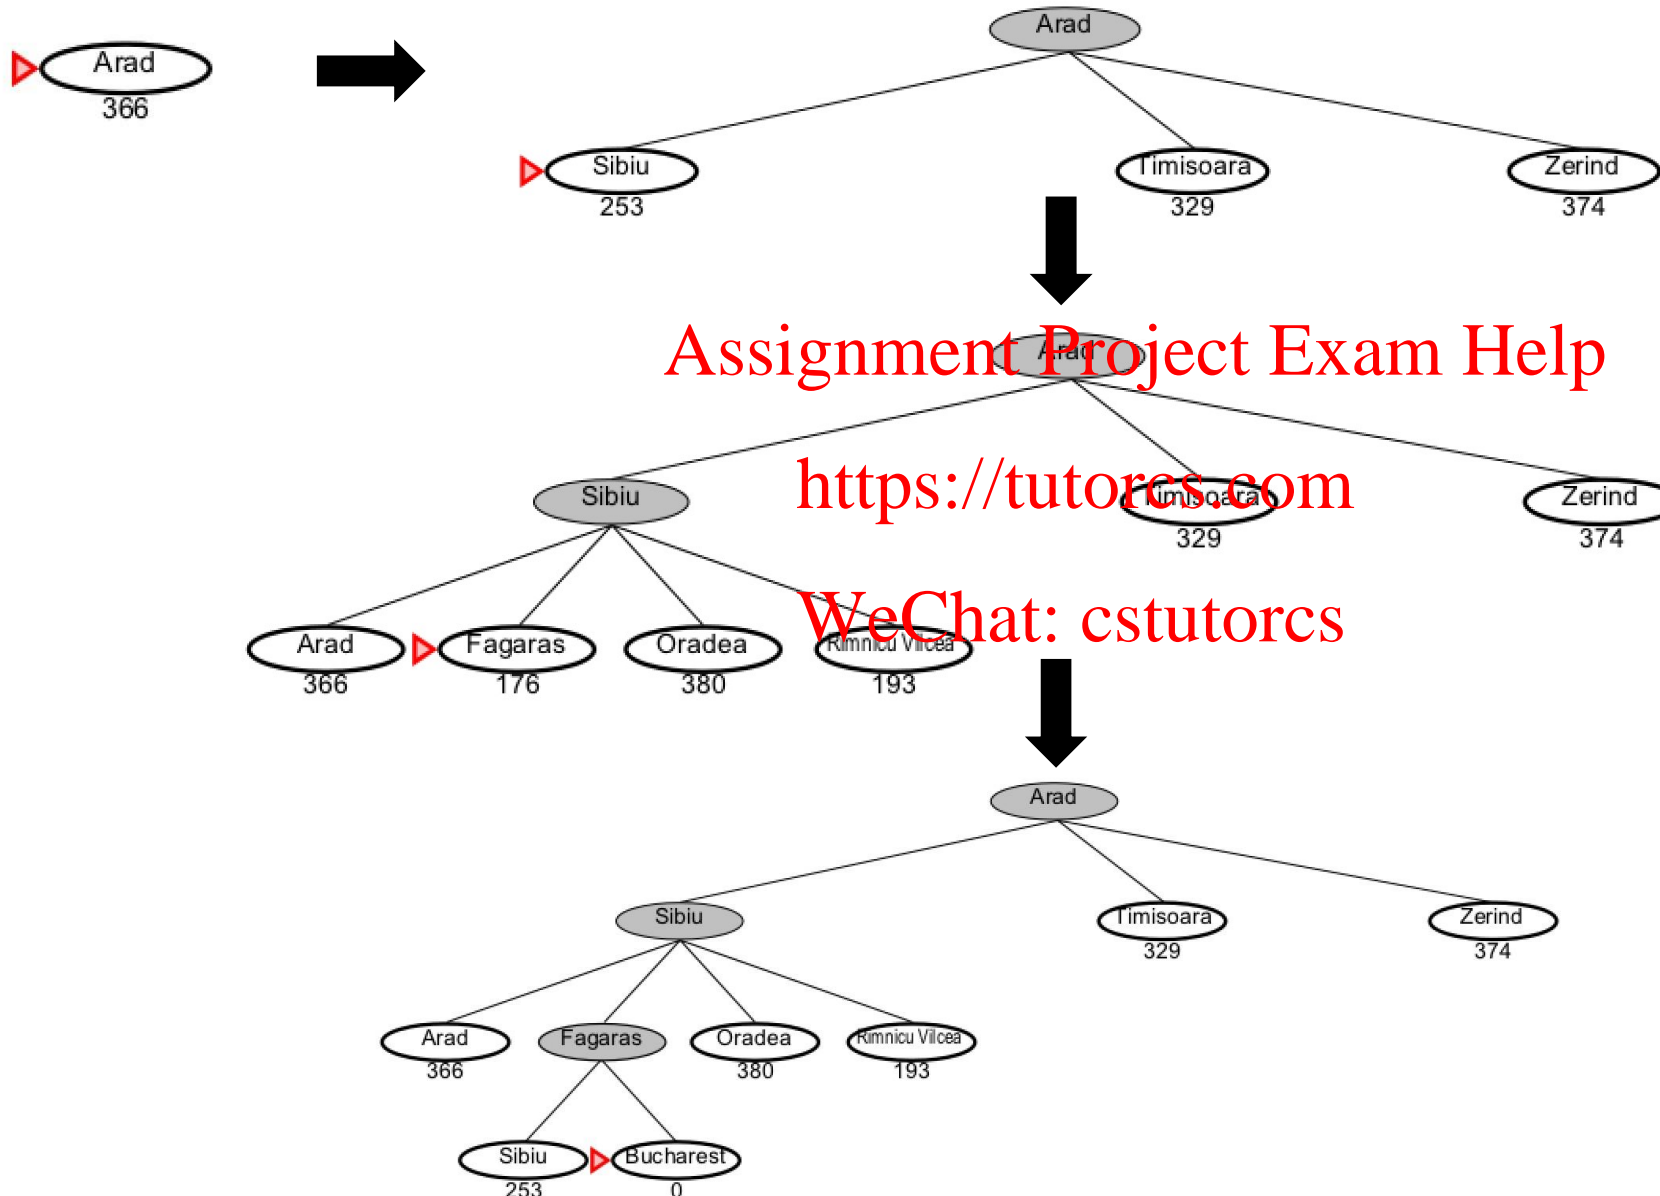


Heuristic function=

Straight-line distance  
to Bucharest

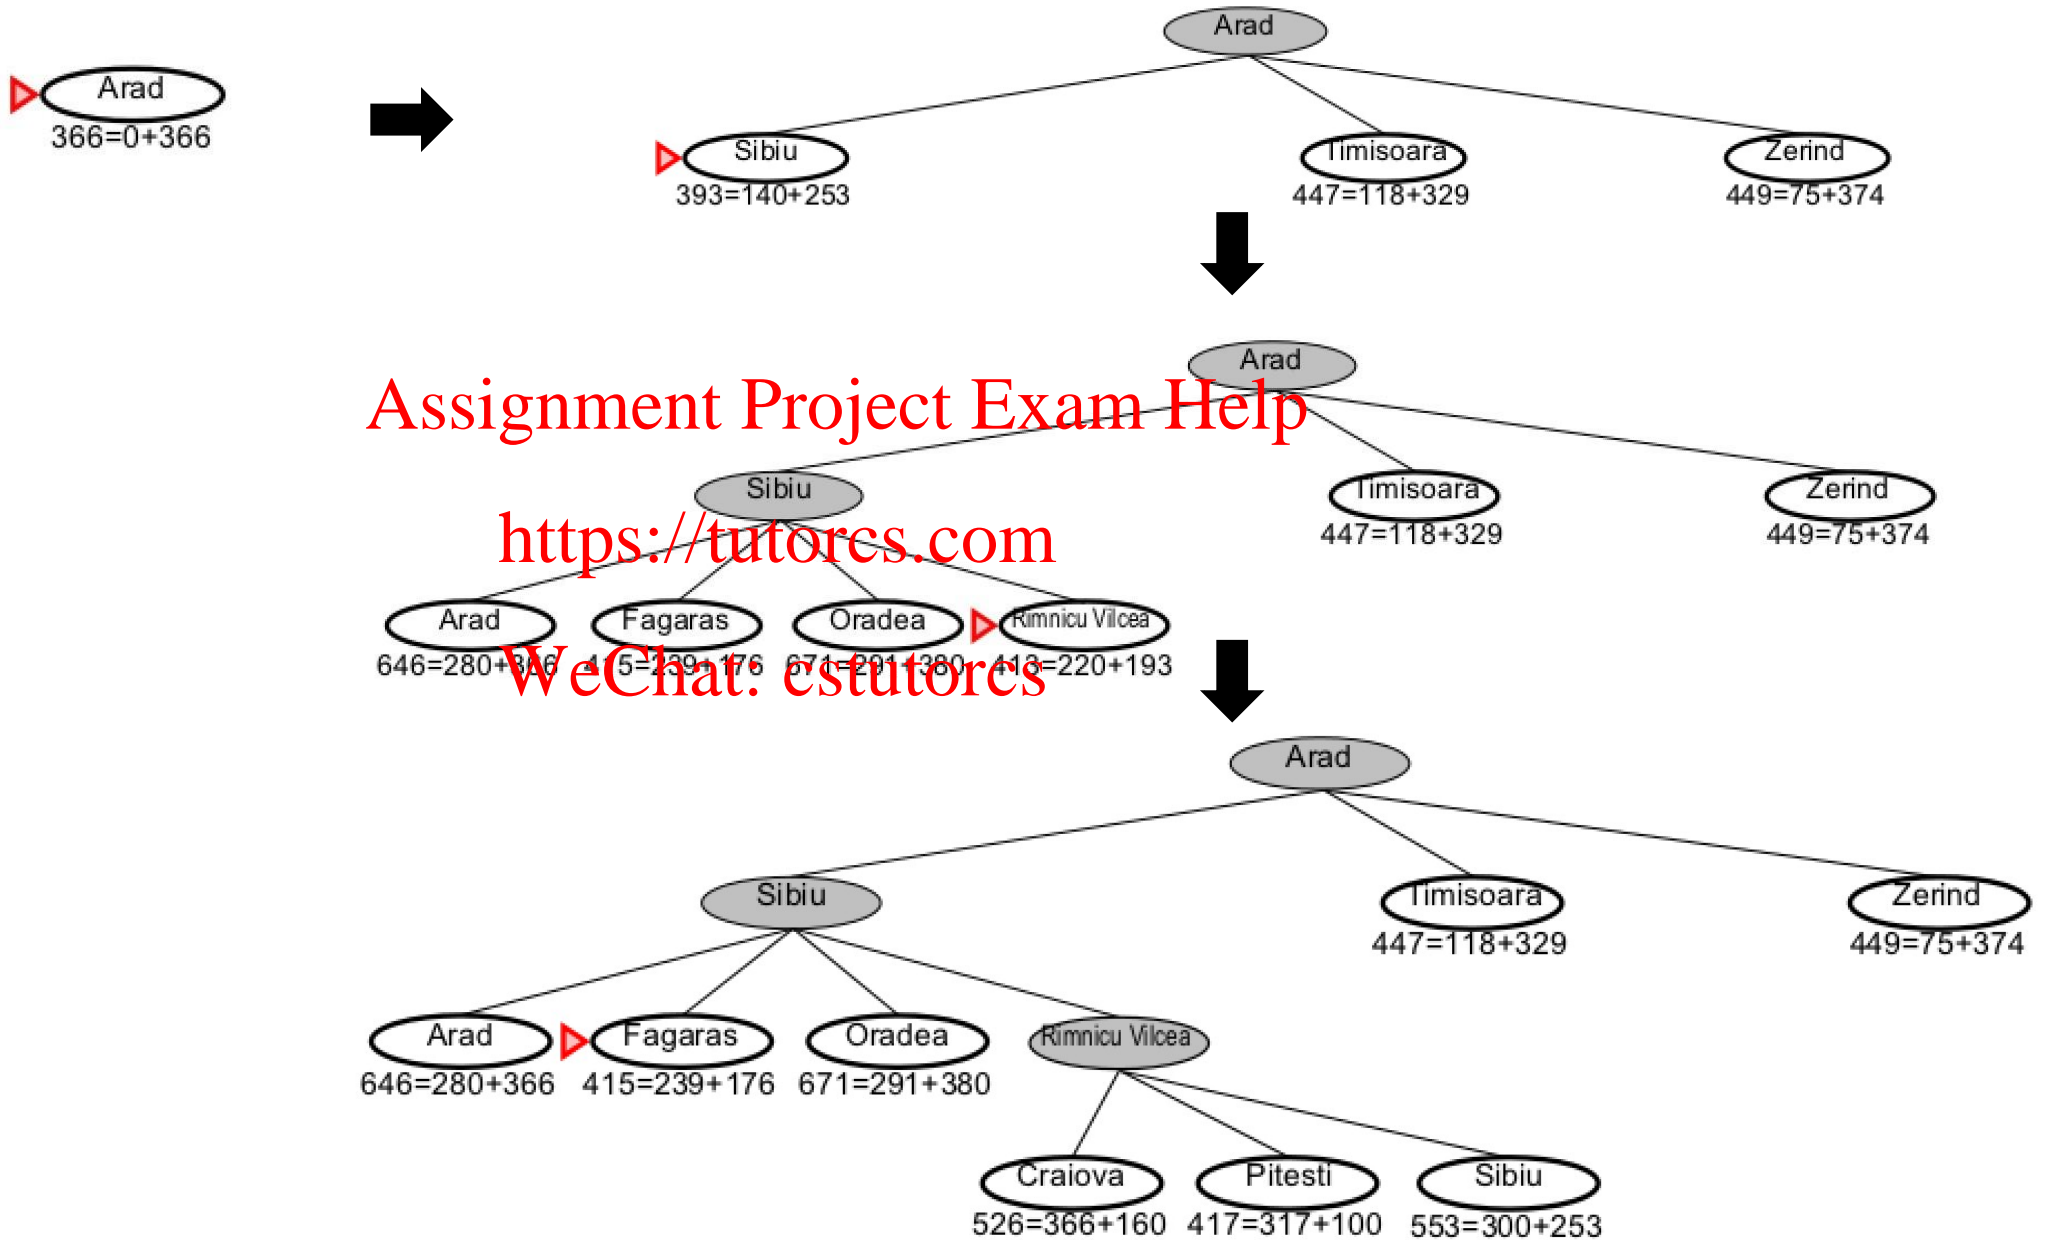
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

- Greedy-best-first (**choose** node with lowest value of heuristic function)



Straight-line distance to Bucharest	
<b>Arad</b>	366
<b>Bucharest</b>	0
<b>Craiova</b>	160
<b>Dobreta</b>	242
<b>Eforie</b>	161
<b>Fagaras</b>	178
<b>Giurgiu</b>	77
<b>Hirsova</b>	151
<b>Iasi</b>	226
<b>Lugoj</b>	244
<b>Mehadia</b>	241
<b>Neamt</b>	234
<b>Oradea</b>	380
<b>Pitesti</b>	98
<b>Rimnicu Vilcea</b>	193
<b>Sibiu</b>	253
<b>Timisoara</b>	329
<b>Urziceni</b>	80
<b>Vaslui</b>	199
<b>Zerind</b>	374

- A\* (choose node with lowest value of actual cost+heuristic function)

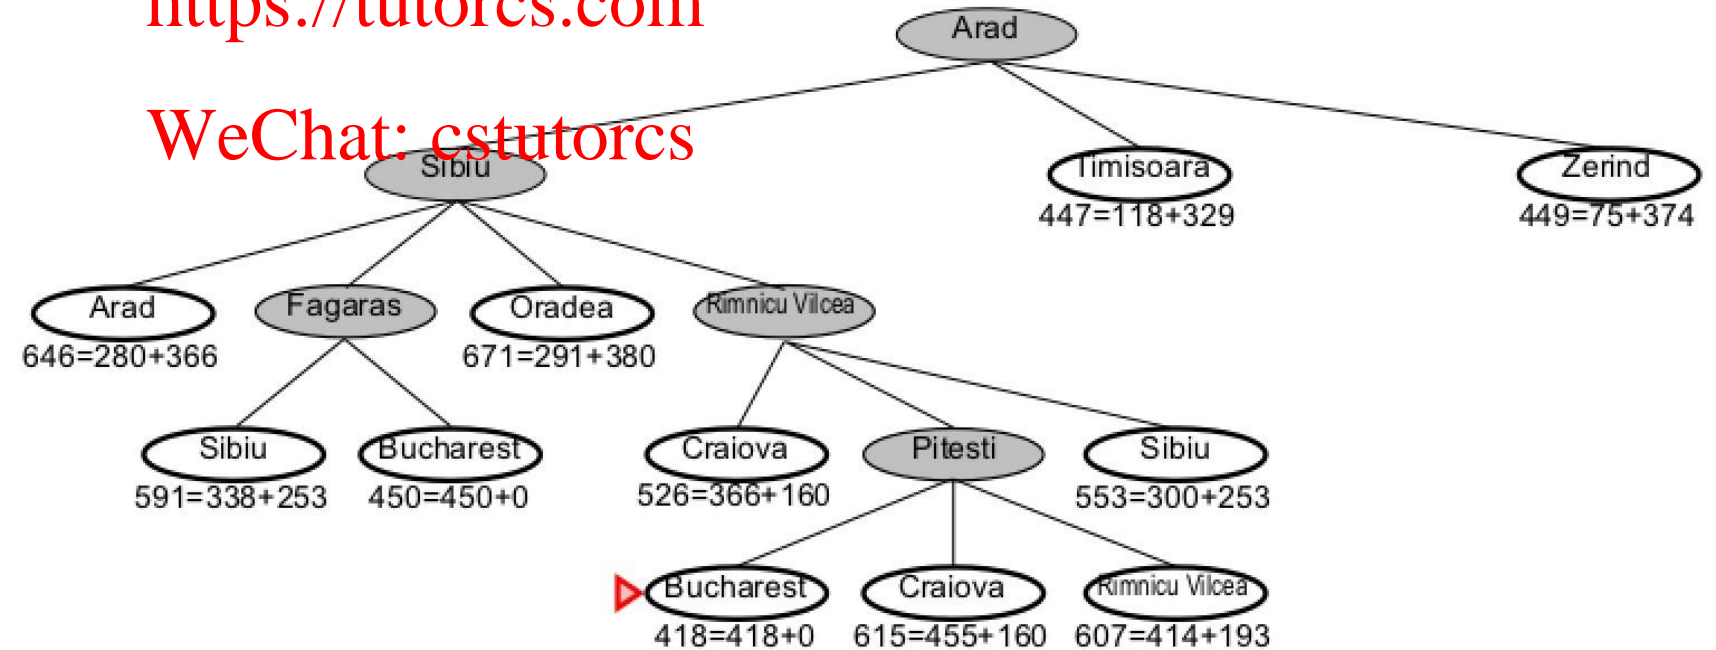




Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



# Properties of heuristic functions

For  $n$  any node, let

$g(n, m)$  be the actual cost of reaching node  $m$  from  $n$

$h(n)$  be the heuristic function applied to  $n$  (assume  $h(n) \geq 0$  -- so that  $h(goal) = 0$ )

- **Admissible:** it never overestimates the actual cost (to goal) – e.g. straight-line distance

$h(n) \leq g(n, m_{goal})$  for  $m_{goal}$  the cheapest goal reachable from  $n$

- **Consistent/Monotonic:** it never overestimates the actual cost to any successor node+the heuristic function applied to that node – e.g. straight-line distance

$h(n) \leq g(n, n') + h(n')$  for  $n'$  any successor of  $n$

Consistent/Monotonic  $\rightarrow$  Admissible

# Properties of informed search (with or **without** loop-checking)

Criterion	Greedy (with)	Greedy ( <b>without</b> )	A* (with)	A* ( <b>without</b> )
Complete?	Yes <sup>†</sup>	No	Yes* <sup>o</sup>	Yes* <sup>o</sup>
Optimal?	No	No	Yes <sup>◇</sup>	Yes <sup>‡</sup>

Assignment Project Exam Help

<https://tutorcs.com>

<sup>†</sup> **s** is finite

\* **b** is finite

<sup>o</sup> step cost is >0

<sup>◇</sup> **h** consistent/monotonic

<sup>‡</sup> **h** admissible

WeChat: cstutorcs

**s**—overall size of the state space

**b**—maximum branching factor of the search tree (may be  $\infty$ )

**m**—maximum depth of the state space (may be  $\infty$ )

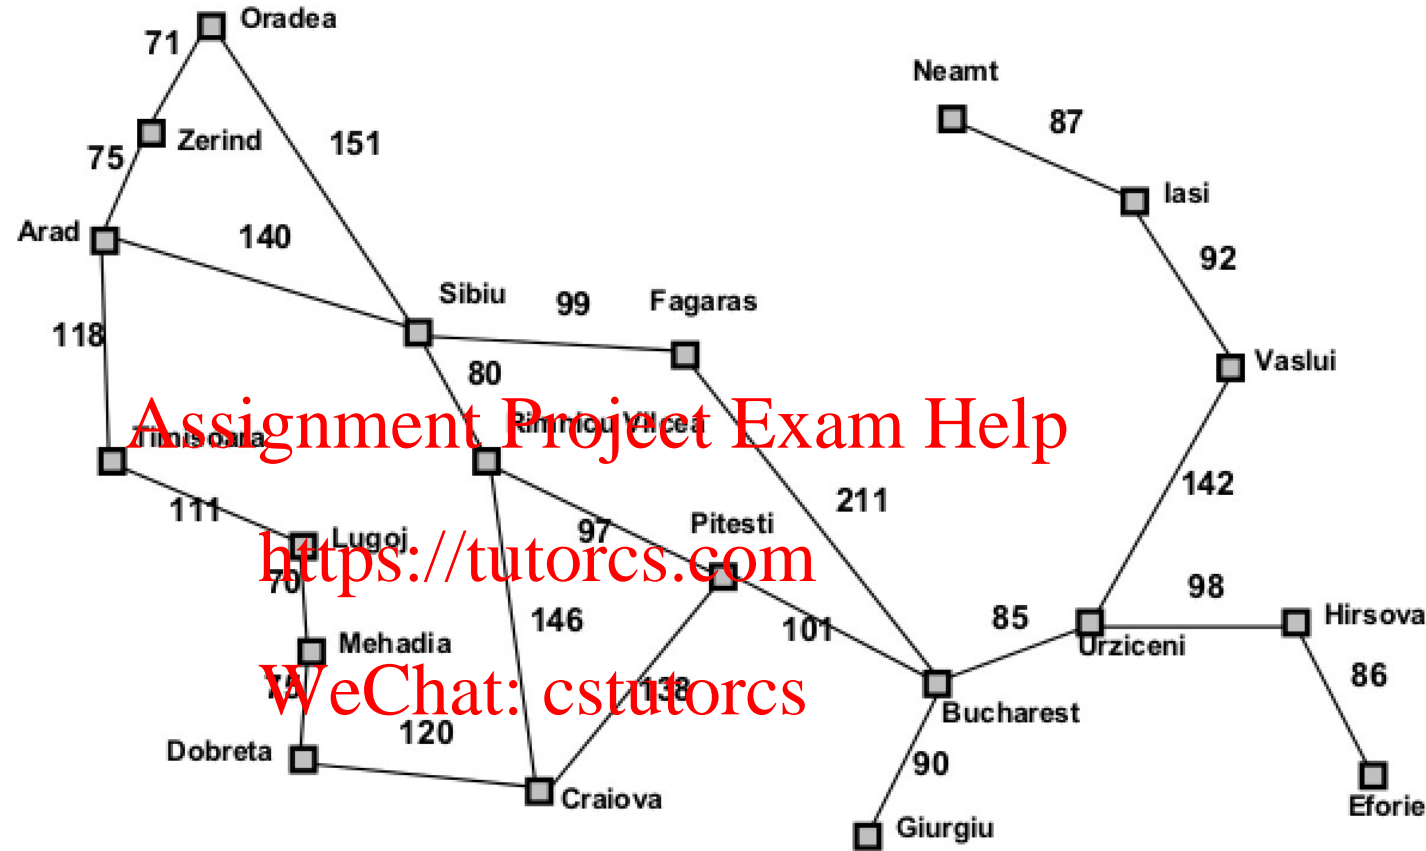
**h**—heuristic function

A\* is **optimally efficient** (no other optimal search algorithm is guaranteed to expand fewer nodes) but it often runs out of **space**



- Greedy best-first search not optimal:

Path via Sibiu-Fagaras is (32 km) longer than path through RV-Pitesti



Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

- Greedy best-first search without loop checking is not complete:  
 Iasi → Neamt → Iasi → Neamt →...



# How to identify **admissible/consistent** heuristics?

- Identify a **relaxed** version of the search problem (with a larger graph – with more edges)  
<https://tutorcs.com>  
WeChat: cstutorcs
- Use cost of optimal solutions for the relaxed problem as (admissible/consistent) heuristics for the original search problem

# 8-puzzle example – relaxed problems

7	2	4
5		6
8	3	1

Start

1	2	3
4	5	6
7	8	

Goal

A tile can move from square A to square B if

- A is horizontally/vertically adjacent to B, and
- B is blank

Assignment Project Exam Help

Relaxed problems:

1. A tile can move from A to B if A is horizontally/vertically adjacent to B
2. A tile can move from A to B if B is blank
3. A tile can move from (any) A to (any) B

Heuristics:

1.  $h(n)$  = Manhattan distance (sum of the distances of tiles as in  $n$  from their goal positions)
2.  $h(n)$  = number of switches (blank-non-blank tile) to obtain goal from  $n$
3.  $h(n)$  = number of misplaced tiles in  $n$  wrt the goal

<https://tutorcs.com>

WeChat: cstutorcs

# Heuristic search – today (example-non-examinable)



Artificial Intelligence  
Assignment Project Exam Help  
Available online 27 December 2017

In Press, Accepted Manuscript – Note to users  
<https://tutorcs.com>



WeChat: cstutorcs

## Star-Topology Decoupled State Space Search

Daniel Gnad , Jörg Hoffmann 

 [Show more](#)

<https://doi.org/10.1016/j.artint.2017.12.004>

[Get rights and content](#)

# Omitted

- Local search
  - Genetic algorithms
  - Simulated annealing
- Non-deterministic , partially observable or unknown environments

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Adversarial search

- Competitive environment
  - opponents with conflicting goals
- “Unpredictable” opponent
  - solution is a strategy (=specifying a move for every possible opponent’s move/sequence of moves)
- Hard to solve, time limits
  - must approximate

# Types of games

Assignment Project Exam Help

perfect information

imperfect information

<u>deterministic</u>	chance
chess, checkers, go, othello	backgammon monopoly
battleships, blind tictactoe	bridge, poker, scrabble nuclear war

<https://tutorcs.com>  
WeChat: cstutorcs

# Adversarial search problems – formally

Two-player, zero-sum (constant-sum) games

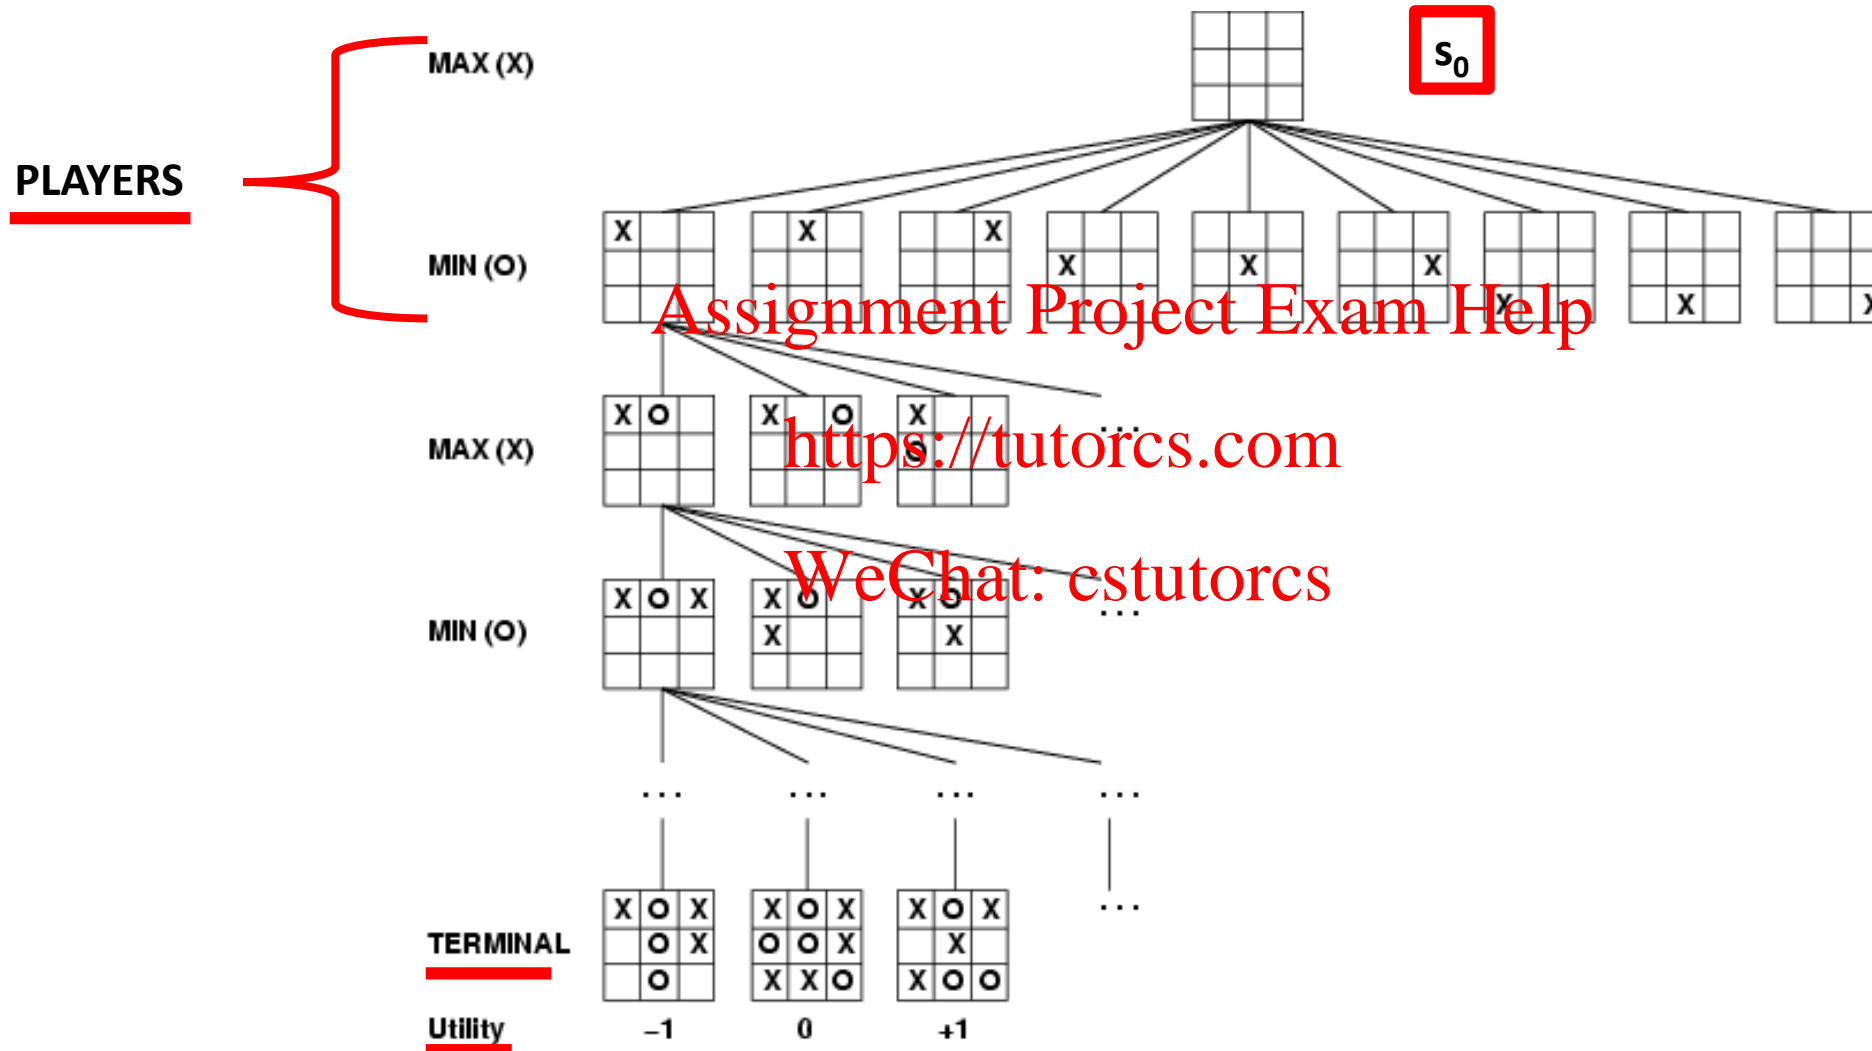
- $s_0$  : initial state
- **PLAYER**( $s$ ): determines which player (of two) moves in a state
- **ACTIONS**( $s$ ): returns the set of legal moves in a state
- **RESULT**( $s,a$ ): returns the state resulting from a move in a state
- **TERMINAL**( $s$ ): is true if the game has ended, false otherwise.
- **UTILITY**( $s,p$ ): returns 1 (win), -1 (lose), 0(draw)  
or 1 (win), 0 (lose), 1/2 (draw)

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Example of two-player, zero-sum game: tic-tac-toe





Perfect play for deterministic, perfect-info games

- Minimax strategy
- $\alpha$ - $\beta$  pruning

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# MINIMAX( $s$ ) – best payoff against best play

for  $\neg \text{TERMINAL}(s)$ :

- if  $\text{PLAYER}(s) = \text{MAX}$ : return  $\max_{\text{actions } a} \text{MINIMAX}(\text{RESULT}(s, a))$
- if  $\text{PLAYER}(s) = \text{MIN}$ : return  $\min_{\text{actions } a} \text{MINIMAX}(\text{RESULT}(s, a))$

for  $\text{TERMINAL}(s)$ :


- return  $\text{UTILITY}(s)$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

e.g. two-ply game



# Properties of minimax

Criterion	minimax
Complete?	Yes*
Time	$O(b^m)$
Space	$O(bm)$
Optimal?	Yes

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

\* If search tree finite

**b**—maximum branching factor of the search tree (may be  $\infty$ )  
**m**—maximum depth of the state space (may be  $\infty$ )

For chess,  $b \approx 35$ ,  $m \approx 100$  for "reasonable" games: not feasible

# $\alpha$ - $\beta$ pruning

- $\alpha$  minimum value that MAX is guaranteed to achieve so far (in some state)
- $\beta$  maximum value that MAX is guaranteed to achieve so far (in some state)

Determine for initial state/node  $s$ :  $\text{max-value}(s, -\infty, +\infty)$  where

Assignment Project Exam Help

$\text{max-value}(s, \alpha, \beta)$ :

<https://tutorcs.com>

WeChat: cstutorcs

- if  $\text{TERMINAL}(s)$  then  
return  $\text{UTILITY}(s)$
- else, starting from  $v = -\infty$ ,
- for each action  $a$ :
  - let  $v = \max(v, \text{min-value}(\text{RESULT}(s, a), \alpha, \beta))$
  - If  $v \geq \beta$  then return  $v$
  - else let  $\alpha = \max(\alpha, v)$

return  $v$

$\text{min-value}(s, \alpha, \beta)$ :

...

$\text{min-value}$

...  $+\infty$

...

...  $\text{min}(v, \text{max-value}(\text{RESULT}(s, a), \alpha, \beta))$

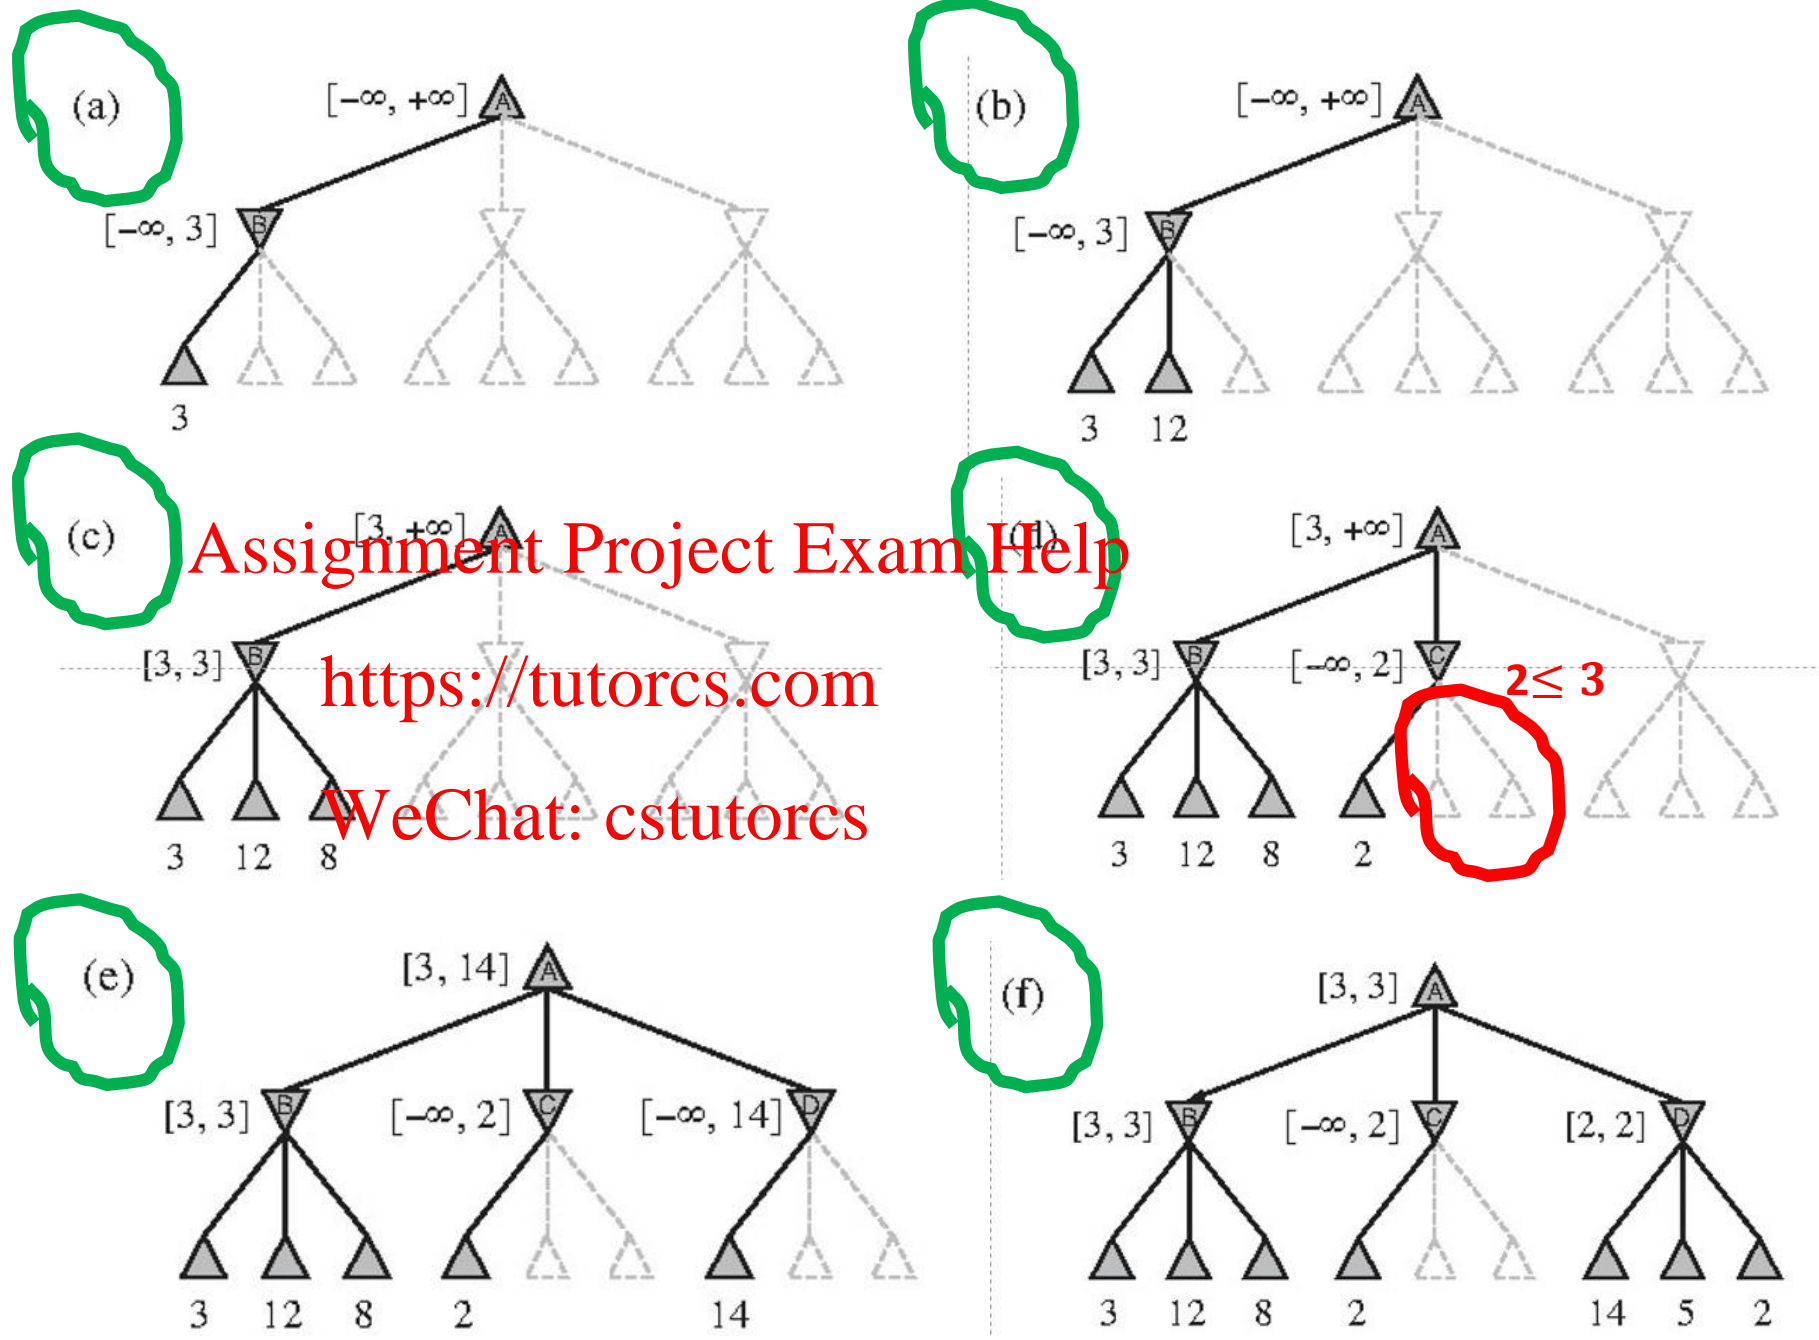
If  $v \leq \alpha$  then return  $v$

else let  $\beta = \min(\beta, v)$

...

$\alpha$ - $\beta$  pruning

$[\alpha, \beta]$



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

# Properties of $\alpha$ - $\beta$ pruning

Criterion	minimax
Complete?	Yes*
Time	$O(b^{m/2})$ $\square$
Space	$O(bm)$
Optimal?	Yes

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

\* If search tree finite

$\square$  if successors are examined best-first

**b**—maximum branching factor of the search tree (may be  $\infty$ )

**m**—maximum depth of the state space (may be  $\infty$ )

For chess,  $b \approx 35$ ,  $m \approx 100$  for "reasonable" games: still slow

# Omitted

## Resource limits:

- Cutoff-Test instead of TERMINAL
- Evaluation instead of UTILITY
  - for chess weighted sum of features (e.g. number of white queens-number of black queens....)

## Other types of games

	deterministic	<u>chance</u>
perfect information	chess, checkers, go, othello	backgammon monopoly
<u>imperfect information</u>	battleships, blind tictactoe	bridge, poker, scrabble nuclear war

# Search&Games – today (example – non-examinable)



Altmetric: 2194

Citations: 1

[More detail >](#)

Assignment Project Exam Help

Article

<https://tutorcs.com>

## Mastering the game of Go without human knowledge

WeChat: cstutorcs

David Silver✉, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel & Demis Hassabis

*Nature* **550**, 354–359 (19 October 2017)

doi:10.1038/nature24270

[Download Citation](#)

Received: 07 April 2017

Accepted: 13 September 2017

Published online: 18 October 2017



# Search - summary

- Formulation of search problems
- Uninformed search algorithms: breadth-first, uniform-cost, depth-first, limited-depth, iterative deepening, backtracking, bidirectional
- Informed (heuristic-based) search algorithms: greedy-best-first, A\*, admissible and consistent/monotonic heuristic functions
- Adversarial search: minimax,  $\alpha$ - $\beta$  pruning
- Properties of search algorithms: completeness, optimality, time and space complexity

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs