

# Introduction to AI

Assignment Project Exam Help

## Non-monotonic Reasoning

<https://tutorcs.com>

WeChat: cstutorcs

Francesca Toni

(thanks to Marek Sergot)


# Outline

- Classical logic: the qualification problem
- Closed World Assumption
- (non-)Monotonicity
- Non-monotonic – defeasible – reasoning
- Negation-as-failure in logic programming (and Prolog)

Assignment Project Exam Help


<https://tutorcs.com>

Recommended:  
Section 5.6



+ sections 1-3 of paper  
on stable models (on cate)

Additional :  
Section 9.4.5  
Section 12.6



We will use the Prolog convention on variables/constant/function symbols. Clauses/rules will be implicitly universally quantified.

# The Qualification Problem (1)

“All birds can fly...”

$\text{flies}(X) \leftarrow \text{bird}(X)$

Assignment Project Exam Help

“...unless they are penguins...”

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X)$

<https://tutorcs.com>

“...or ostriches...” WeChat: cstutorcs

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X), \neg \text{ostrich}(X)$

“...or wounded...”

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X), \neg \text{ostrich}(X), \neg \text{wounded}(X)$

“...or dead, or sick, or glued to the ground, or...”

# The Qualification Problem (2)

Let **BIRDS** be the set of sentences about flying birds.

Even if we could list all these exceptions, classical logic would still not allow:

Assignment Project Exam Help  
 $\text{BIRDS} \cup \{\text{bird}(\text{tweety})\} \not\models \text{flies}(\text{tweety})$   
<https://tutorcs.com>

unless we also affirm all the *qualifications*, viz.:

- penguin(tweety)
- wounded(tweety)
- sick(tweety)
- ostrich(tweety)
- dead(tweety)
- glued\_to\_the\_ground(tweety)

...

# Classical logic is inadequate

What we want is a new kind of ‘entailment’:

$BIRDS \cup \{bird(tweety)\} \models^* flies(tweety)$

Namely, from BIRDS and  $bird(tweety)$  it follows *by default* – in the absence of information to the contrary – that  $flies(tweety)$ .

This kind of reasoning will be *defeasible/non-monotonic*.

# Non-monotonic logics

Classical logic is *monotonic*. For a set of sentences S:

if  $S \models \alpha$  then  $S \cup X \models \alpha$

Namely, new information X always preserves old conclusions  $\alpha$ .

**Assignment Project Exam Help**

Reasoning by default is typically *non-monotonic*. We may have that:

~~<https://tutorcs.com>~~

Examples:

- $\text{DRINKS} \cup \{\text{coffee}\} \models^* \text{tastes nice}$
- $\text{DRINKS} \cup \{\text{coffee}\} \cup \{\text{diesel oil}\} \not\models^* \text{tastes nice}$
- $\text{BIRDS} \cup \{\text{bird(tweety)}\} \models^* \text{flies(tweety)}$
- $\text{BIRDS} \cup \{\text{bird(tweety)}\} \cup \{\text{penguin(tweety)}\} \not\models^* \text{flies(tweety)}$

There is a branch of AI focusing on non-monotonic logic for default reasoning – within knowledge representation and reasoning in AI.

# The 'Closed World Assumption'

Consider the set of sentences that could be describing a database DB:

has-office-in(IBM, Winchester)      city(Winchester)  
has-office-in(IBM, London)      city(London)  
has-office-in(IBM, Paris)      city(Paris)  
has-office-in(MBI, London)      city(NewYork)  
capital-city(London)      company(IBM)  
capital-city(Paris)      company(MBI)

WeChat: cstutorcs

Does MBI have an office in Paris?

Does IBM have an office in New York?

*We do not know.*

But it is usual in databases (and many other contexts) to make a *Closed World Assumption* (CWA):

if  $\alpha$  is not in the DB, assume  $\neg\alpha$

# Normal logic programs

A normal logic program is a set of clauses of the form:

$$A \leftarrow L_1, \dots, L_n \quad (n \geq 0)$$

where  $A$  is an atom and each  $L_i$  is a literal.

A literal is either an atom (a 'positive literal') or of the form  
 $\text{not } B$

where  $B$  is an atom. ( $\text{not } B$  is a 'negative literal').

The atom  $A$  is the head of the clause; the literals  $L_1, \dots, L_n$  are the body of the clause. When the body is empty ( $n = 0$  above) the arrow  $\leftarrow$  is usually omitted.

$\text{not}$  is negation as failure (NAF): informally (for now)


$\text{not } B$  succeeds when all attempts to prove  $B$  fail



# SLD+Negation-as-failure: SLDNF (1)

We have already seen the computation of a goal (query)

$G = L_1, \dots, L_m$  as a series of derivation steps:



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- $\theta_i$  is the mgus at the  $i$ -th derivation step
- The answer computed is (the restriction to the vars of  $G$  of) the composition of all these mgus:

$$\theta = \theta_1 \circ \dots \circ \theta_n$$

# SLD+Negation-as-failure: SLDNF (2)

Now there are two kinds of derivation steps:

(a) select a positive literal  $L_j = B$  from the current goal:

$\leftarrow L_1, \dots, L_{j-1}, B, L_{j+1}, \dots, L_n$   
 $\theta_i \mid$  match  $B$  with  $B' \leftarrow M_1, \dots, M_k$ , with  $B\theta_i = B'\theta_i$   
 $\leftarrow (L_1, \dots, L_{j-1}, M_1, \dots, M_k, L_{j+1}, \dots, L_n)\theta_i$

<https://tutorcs.com>

(b) select a negative literal  $L_j = \text{not } B$  from the current goal:

$\leftarrow L_1, \dots, L_{j-1}, \text{not } B, L_{j+1}, \dots, L_n$   
 $\theta_i = \{\}$  subcomputation:  
 all ways of computing goal  $B$  must fail (finitely)  
 $\leftarrow (L_1, \dots, L_{j-1}, L_{j+1}, \dots, L_n)$

Note that the NAF sub-computation just **checks not B**: it does not generate bindings (substitutions) for variables.

# SLDNF Example (S is ground)

S:  $p \leftarrow q, \text{not } r$

$r \leftarrow b$

$r \leftarrow \text{not } q$


$q \leftarrow$

Goal:  $p$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



computed answer:  $\{\}$

# SLDNF Example (S is not ground)

S:  $p(X) \leftarrow q(X), \text{not } r(X)$

$r(X) \leftarrow b(X)$

$r(X) \leftarrow \text{not } q(X)$


$q(a) \leftarrow$

Assignment Project Exam Help

Goal: p

<https://tutorcs.com>

WeChat: cstutorcs



computed answer:  $\{X/a\}$

# Selection of sub-goals in SLDNF

- Sub-goals can be selected in any order. The answers are the same, whatever the selection rule
  - Prolog selects always the leftmost sub-goal in the current goal.
- However, the selection of sub-goals must be **safe**: it must not pick a non-ground negative literal
  - Most Prolog systems do not implement this!

E.g.  $S = \{p(a) \leftarrow, q(b) \leftarrow\}$

$\leftarrow \text{not } p(X), q(X)$

...

fail

$\leftarrow \text{not } p(X), \underline{q(X)}$

$\leftarrow \text{not } p(b)$

...  
□

Wrongly interprets  $\exists X \neg p(X)$  as  $\neg \exists X p(X)$

# SLDNF is non-monotonic

Consider the set of sentences S:

$p(X) \leftarrow q(X), \text{ not } r(X)$

$q(a)$  Assignment Project Exam Help

$r(b)$

Then  $S \vdash_{\text{SLDNF}} p(a)$

but  $S \cup \{r(a)\} \not\vdash_{\text{SLDNF}} p(a)$

<https://tutorcs.com>

WhatsApp: [cstutorcs](https://tutorcs.com)

Note: SLDNF builds in a kind of Closed World Assumption:

$S \cup \{r(a)\} \vdash_{\text{SLDNF}} \neg p(a)$

# Solving the qualification problem using NAF

## Example

- Typically (by default, unless there is reason to think otherwise,...) a bird can fly.
- Except that dead birds cannot fly

<https://tutorcs.com>  
 $\text{can\_fly}(X) \leftarrow \text{bird}(X), \text{not abnormal\_bird}(X)$   
 $\text{abnormal\_bird}(X) \leftarrow \text{dead}(X)$

## Example

People are innocent by default (unless they can be proven to be guilty)

$\text{innocent}(X) \leftarrow \text{not guilty}(X)$

# Credulous vs Sceptical non-monotonic reasoning: “The Nixon diamond”

- Quakers are typically pacifists.
- Republicans are typically not pacifists.
- Richard Nixon is a Quaker
- Richard Nixon is a Republican

Assignment Project Exam Help

Do we conclude that Nixon is a pacifist or not?

<https://tutores.com>

*Sceptical (or cautious) reasoning:*

WeChat: ctutores

No, since we cannot choose between two possible, conflicting default conclusions.

*Credulous (or brave) reasoning:*

Yes, to both, since we have reason to believe both.

Which form of reasoning to choose? It depends on the needs of our application/problem



# The Nixon diamond using NAF

- Quakers are typically pacifists.
- Republicans are typically not pacifists.
- Richard Nixon is a Quaker
- Richard Nixon is a Republican

Do we conclude that Nixon is a pacifist or not?

<https://tutorcs.com>  
pacifist(X) ← quacker(X), not abnormal\_quacker(X)  
not\_pacifist(X) ← republican(X), not abnormal\_republican(X)

abnormal\_quacker(X) ← not\_pacifist(X)  
abnormal\_republican(X) ← pacifist(X)

quacker(nixon)  
republican(nixon)

# SLDNF for the Nixon diamond

By shortening predicates (pacifist becomes p etc)

$p(X) \leftarrow q(X), \text{ not } ab\_q(X)$

$n\_p(X) \leftarrow r(X), \text{ not } ab\_r(X)$

$ab\_q(X) \leftarrow n\_p(X)$

$ab\_r(X) \leftarrow p(X)$

$q(\text{nixon})$

$r(\text{nixon})$

Assignment Project Exam Help

Goal (for example. Here nixon becomes nix):  $p(\text{nix})$   
<https://tutorcs.com>

$\leftarrow p(\text{nix})$

$\leftarrow q(\text{nix}), \text{ not } ab\_q(\text{nix})$

$\leftarrow \text{not } ab\_q(\text{nix})$



$\leftarrow ab\_q(\text{nix})$

$\leftarrow n\_p(\text{nix})$

$\leftarrow r(\text{nix}), \text{ not } ab\_r(\text{nix})$

$\leftarrow \text{not } ab\_r(\text{nix})$



$\leftarrow ab\_r(\text{nix})$

...

**Infinite computation!**

# Semantics for normal logic programs

- Completion semantics
  - Well-founded model \*
  - Stable models (see answer sets in Part II) \*
- <https://tutorcs.com>

WeChat: cstutorcs

\* these semantics equate each **normal logic program**  $S$  to the set of all its ground instances over the underlying Herbrand universe

# Semantics of sets of definite clauses = positive logic programs (recap)

Given a set of definite clauses  $S$ :

- the *Herbrand universe* is the set of all ground terms obtained from constants and function symbols in  $S$ , and the *Herbrand base* ( $HB$ ) is the set of all atoms from predicate symbols in  $S$  and terms in the Herbrand universe
- A *Herbrand model* is a subset of  $HB$  that renders  $S$  true
- The meaning of  $S$  is the *least Herbrand model* ( $LHM$ ) of  $S$

Note: the LHM is the least fixed point of the (continuous) immediate consequence operator

# From positive to normal logic program semantics

Given a normal logic program  $S$ :

$$T_S(X) = \{a \in HB \mid a \leftarrow b_1, \dots, b_m, \text{not } b_{m+1}, \dots, \text{not } b_{m+n} \in S, \\ \{b_1, \dots, b_m\} \subseteq X, \\ \{b_{m+1}, \dots, b_{m+n}\} \cap X = \{\}\}$$

WeChat: cstutorcs

- $T_S$  is no longer guaranteed to be continuous, as it may not be monotonic:
  - E.g.  $S = \{p \leftarrow \text{not } q\}$   
 $T_S(\{\}) = \{p\}$   
 $T_S(\{q\}) = \{\}$

# Completion semantics (1)

1. Rewrite all rules in S as rules of the form (with  $\wedge$ ,  $\neg$ )

$$p(\vec{X}) \leftarrow \vec{X} = \vec{t}, \text{ body}$$

– e.g.  $p(X,3) \leftarrow q(X)$ , not  $r(3)$

becomes  $p(X,Y) \leftarrow Y=3 \wedge q(X) \wedge \neg r(3)$

Assignment Project Exam Help

2. Combine all rules with the same head

$$p(\vec{X}) \leftarrow \vec{X} = \vec{t}_1 \wedge \text{body}_1 \dots p(\vec{X}) \leftarrow \vec{X} = \vec{t}_n \wedge \text{body}_n$$

<https://tutorcs.com>

$$\text{into } (\forall) p(\vec{X}) \leftrightarrow (\vec{X} = \vec{t}_1 \wedge \text{body}_1) \vee \dots \vee (\vec{X} = \vec{t}_n \wedge \text{body}_n)$$

WeChat: cstutorcs

– e.g.  $p(X,Y) \leftarrow Y=3 \wedge q(X) \wedge \neg r(3)$ ,  $p(X,Y) \leftarrow X=2$

becomes  $(\forall) p(X,Y) \leftrightarrow (Y=3 \wedge q(X) \wedge \neg r(3)) \vee (X=2)$

3. If some  $q(\_)$  (in vocabulary of S) is not the head of any rule in S, add

$$(\forall) q(\vec{X}) \leftrightarrow \text{false}$$

# Completion semantics (2)

4. Finally, add Clark's Equality Theory (CET):

- $f(\vec{t}) \neq g(\vec{s})$  for all pairs of distinct terms in the Herbrand universe of  $S$
- $(\forall) X=X, X=Y \rightarrow Y=X, X=Y \wedge Y=Z \rightarrow X=Z$
- $(\forall) \vec{X} = \vec{Y} \rightarrow p(\vec{X}) = p(\vec{Y})$  for all predicate symbols  $p$  in  $S$
- $\vec{X} \neq t[\vec{X}]$  for all terms  $t$  where  $\vec{X}$  occurs

The resulting logical theory is called  
the **completion** of  $S$

# Completion: example

S:  $p(X) \leftarrow q(X), \text{ not } r(X)$

$r(X) \leftarrow b(X)$

$r(X) \leftarrow \text{not } q(X)$

$q(a) \leftarrow$

Assignment Project Exam Help

<https://tutorcs.com>

Completion of S: CET+

WeChat: cstutorcs

$\forall X (p(X) \leftrightarrow q(X) \wedge \neg r(X))$

$\forall X (r(X) \leftrightarrow b(X) \vee \neg q(X))$

$\forall X (q(X) \leftrightarrow X=a)$

$\forall X (b(X) \leftrightarrow \text{false})$



# Soundness of SLDNF (with respect to the completion semantics)

Let  $\text{Comp}(S)$  be the completion of  $S$

- If there exists an SLDNF-refutation of  $P$  from  $S$  with answer  $\theta$  then  $\text{Comp}(S) \models (\forall) P\theta^*$

<https://tutorcs.com>

WeChat: cstutorcs \* Where not is interpreted as  $\neg$

- If there exists a **finitely failed** SLDNF-refutation of  $P$  from  $S$  then  $\text{Comp}(S) \models \neg \exists P$

Negation as **finite failure**

# Completion semantics for SLDNF: issues

What if SLDNF does not terminate? What does non-termination “mean”?

Assignment Project Exam Help

- $S = \{p \leftarrow q, p \leftarrow \text{not } q, q \leftarrow q\}$ ,  $P = p$

<https://tutorcs.com>

Completion of  $S = \{p \leftrightarrow q \vee \neg q, q \leftrightarrow q\} + \text{CET} \models p$

WeChat: cstutorcs

- $S = \{p \leftarrow \text{not } p\}$ ,  $P = p$

Completion of  $S = \{p \leftrightarrow \neg p\} + \text{CET} \models p, \neg p$   
(inconsistent)

# Well-founded model semantics

## Negation as **infinite failure**

- $S = \{p \leftarrow q, p \leftarrow \text{not } q, q \leftarrow q\}, P = p$

Well-founded model of  $S$  is  $(\{p\}, \{\text{not } q\})$ :

$S$  “wfm-entails”  $P$

- $S = \{p \leftarrow \text{not } p\}, P = p$

Well-founded model of  $S$  is  $(\{\}, \{\})$ :

$S$  “does not wfm-entail”  $P$

SLDNF+tabling: XSB Prolog: <http://xsb.sourceforge.net/>

# Well-founded model

(In, Out) such that

–  $In \subseteq HB$

–  $Out \subseteq HB^{\text{not}} = \{\text{not } a \mid a \in HB\}$

Note:

<https://tutorcs.com>

WeChat: cstutorcs

- atoms in  $HB - (In \cup \{\text{not } a \mid \text{not } a \in Out\})$  are “undecided”
- e.g. for  $S = \{p \leftarrow \text{not } p\}$ , the well-founded model is  $(\{\}, \{\})$  and  $p$  is “undecided”

(definition of well-founded model omitted  
and thus non examinable)

# Nixon diamond using the completion semantics and the well-founded model

$p(V) \leftarrow q(V), \text{ not } ab\_q(V)$   
 $ab\_q(V) \leftarrow n\_p(V)$   
 $q(nix)$

$n\_p(V) \leftarrow r(V), \text{ not } ab\_r(V)$   
 $ab\_r(V) \leftarrow p(V)$   
 $r(nix)$

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- Completion= CEF
 

$p(V) \leftrightarrow q(V) \wedge \neg ab\_q(V)$ $ab\_q(V) \leftrightarrow n\_p(V)$ $q(V) \leftrightarrow V=nix$	$n\_p(V) \leftrightarrow r(V) \wedge \neg ab\_r(V)$ $ab\_r(V) \leftrightarrow p(V)$ $r(V) \leftrightarrow V=nix$
--	--

  - $p(nix), n\_p(nix)$  hold in different models
- Well-founded model =  $(\{q(nix), r(nix)\}, \{\})$ 
  - $p(nix), n\_p(nix)$  are both undecided

# Stable models

- Given a (ground) normal logic program  $S$  and  $X \subseteq \text{HB}$ , the *reduct of  $S$  by  $X$*  (referred to as  $S^X$ ) is obtained in two steps:

1. Eliminate all rules with  $\text{not } p$  in the body, for every  $p \in X$
2. Eliminate all negative literals from the body of all remaining rules

( $S^X$  is a set of definite clauses)

- $X \subseteq \text{HB}$  is a *stable model* of  $S$  iff the LHM of  $S^X$  is  $X$

Several efficient answer set solvers exist – see Part II

# Nixon diamond using stable models (1)

$p(V) \leftarrow q(V), \text{ not } ab\_q(V)$

$n\_p(V) \leftarrow r(V), \text{ not } ab\_r(V)$

$ab\_q(V) \leftarrow n\_p(V)$

$ab\_r(V) \leftarrow p(V)$

$q(nix)$

$r(nix)$

## Assignment Project Exam Help

Here  $HB = \{q(nix), r(nix), p(nix), \text{not } p(nix), ab\_q(nix), ab\_r(nix)\}$ . Let  $S$  be the set of all ground instances of these rules over  $HB$ .

Consider  $X = \{p(nix), ab\_r(nix), q(nix), r(nix)\}$ . **Is  $X$  a stable model? YES.**

To see this, let us construct  $S^X$

1. Eliminates  $n\_p(nix) \leftarrow r(nix), \text{ not } ab\_r(nix)$

2. Gives  $p(nix) \leftarrow q(nix),$

$ab\_q(nix) \leftarrow n\_p(nix), \quad ab\_r(nix) \leftarrow p(nix),$

$q(nix), \quad r(nix)$

The LHM of  $S^X$  is  $\{q(nix), r(nix), p(nix), ab\_r(nix)\} = X$ , as required.

## Nixon diamond using stable models (2)

$X = \{q(\text{nix}), r(\text{nix}), n\_p(\text{nix}), ab\_q(\text{nix})\}$  is also a stable model (Exercise: verify this)

Assignment Project Exam Help

- Multiple “extensions” = stable models
- Credulous (whatever holds in some stable model) vs sceptical (whatever holds in all stable models)

<https://tutorcs.com>

WeChat: cstutorcs



# Summary

- Classical logic: the qualification problem
- Closed World Assumption
- (non-)Monotonicity
- Non-monotonic/defeasible reasoning
- Negation-as-failure/SLDNF for non-monotonic reasoning
- Completion semantics, well-founded semantics (just a mention) and stable model semantics

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs